

Computer Science Tripos - Part II Project

An accelerated, network-assisted TCP recovery

May 17, 2019

Proforma

Name:	Thanh Bui
College:	Downing College
Project Title:	An accelerated, network-assisted TCP recovery
Examination:	Computer Science Tripos – Part II, June 2019
Word Count:	1587¹
Project Originator:	Dr Noa Zilberman
Supervisor:	Dr Noa Zilberman

Original Aims of the Project

To write a demonstration dissertation² using \LaTeX to save student's time when writing their own dissertations. The dissertation should illustrate how to use the more common \LaTeX constructs. It should include pictures and diagrams to show how these can be incorporated into the dissertation. It should contain the entire \LaTeX source of the dissertation and the makefile. It should explain how to construct an MSDOS disk of the dissertation in Postscript format that can be used by the book shop for printing, and, finally, it should have the prescribed layout and format of a diploma dissertation.

Work Completed

All that has been completed appears in this dissertation.

¹This word count was computed by `detex diss.tex | tr -cd '0-9A-Za-z \n' | wc -w`

²A normal footnote without the complication of being in a table.

Special Difficulties

Learning how to incorporate encapsulated postscript into a \LaTeX document on both Ubuntu Linux and OS X.

Declaration

I, Thanh Bui of Downing College, being a candidate for Part II of the Computer Science Tripos, hereby declare that this dissertation and the work described in it are my own work, unaided except as may be specified below, and that the dissertation does not contain material that has already been used to any substantial extent for a comparable purpose.

SIGNED

DATE

Contents

List of figures	9
1 INN	11
1.1 Motivation	11
1.1.1 Background	11
1.2 Problem Outline	12
1.3 Related Work	12
2 Preparation	15
2.1 Transmission Control Protocol	15
2.2 Starting Point	15
2.3 Software Used	15
2.3.1 P4-NetFPGA Platform	15
2.3.2 P4 Programming Language	15
2.4 Previous Work	15
3 Implementation	17
3.1 The SimpleSumeSwitch Architecture	17
3.2 The Cache Queue	17
4 Evaluation	19
5 Conclusion	21
5.1 Results	21
5.2 Lessons Learnt	21
5.3 Future Work	21
Bibliography	23

List of figures

1.1	The standard convention of TCP handling.	11
1.2	The proposed TCP handling.	12

Acknowledgements

This document owes much to an earlier version written by Simon Moore [?]. His help, encouragement and advice was greatly appreciated.

Chapter 1

INN

1.1 Motivation

1.1.1 Background

Currently, when there is packet loss, signaled by duplicate ACKs (DUP ACKs), the DUP ACKs will traverse all the way back to the host (**Figure 1.1**). The host receives the DUP ACKs, then resends the packet. The goal of my project is to implement a fast TCP recovery mechanism, which aims to reduce the response time to DUP ACKs and reduce changes to the congestion window by retransmitting packets from within the network (e.g., from the switch), instead of sending the DUP ACKs back to the host (**Figure 1.2**). The implementation will be based on KVS concepts, where the key is the flow Id and sequence number, and the value is the packet.

For this project, I will be using the P4 programming language [1]. It is a language designed to allow the programming of packet forwarding planes. Besides, unlike general purpose languages such as C or Python, P4 is domain-specific with a number of constructs optimized around network data forwarding, hence is well suited to such a network application.

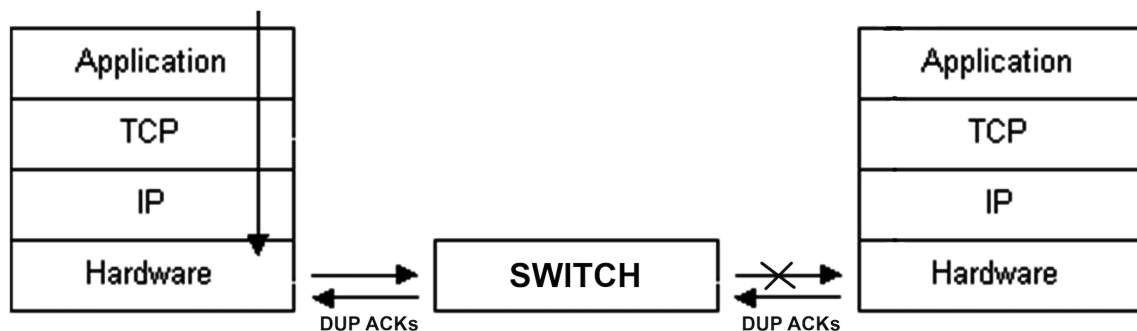


Figure 1.1: The standard convention of TCP handling.

1.2. PROBLEM OUTLINE

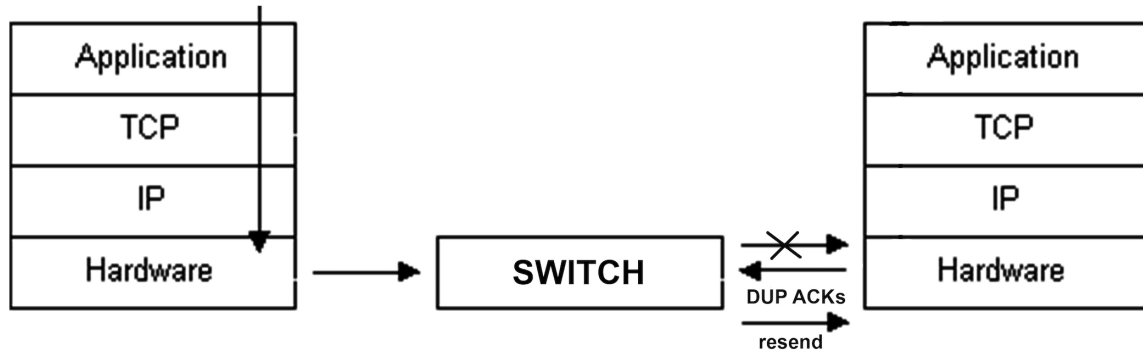


Figure 1.2: The proposed TCP handling.

1.2 Problem Outline

1.3 Related Work

Test

Chapter 2

Preparation

a

2.1 Transmission Control Protocol

2.2 Starting Point

2.3 Software Used

2.3.1 P4-NetFPGA Platform

a

2.3.2 P4 Programming Language

a

2.4 Previous Work

a

a

a

a

a

2.4. PREVIOUS WORK

a

a

a a a

Chapter 3

Implementation

3.1 The SimpleSumeSwitch Architecture

This project will work mainly with a NetFPGA SUME board [2], using P4 programming language. I will be using the P4-NetFPGA workflow, which provides infrastructure to compile P4 programs to NetFPGA [3]. Apart from that, everything else will be built from scratch.

I have no prior experience with either NetFPGA or P4, but this will be mitigated through self-learning in which I will make use of the online tutorials, Google’s resources and the P4 community documentation, as well as the experience of my supervisors.

3.2 The Cache Queue

The relevant Tripos courses that can serve as a starting point for this project are primarily: *Computer Networking*, *Principles of Communications* and *ECAD and Architecture Practical Classes*. Since the courses are introductory, I will also consult Part III’s *High Performance Networking* course. I also plan to bridge any knowledge gap through extensive personal reading as well as help from my project supervisor.

Chapter 4

Evaluation

Tests

Chapter 5

Conclusion

5.1 Results

5.2 Lessons Learnt

5.3 Future Work

Bibliography

- [1] “The P4 Language Consortium,” <https://p4.org/>.
- [2] N. Zilberman, Y. Audzevich, G. A. Covington, and A. W. Moore, “NetFPGA SUME: Toward 100 Gbps as research commodity,” *IEEE micro*, vol. 34, no. 5, pp. 32–41, 2014.
- [3] “NetFPGA/P4-NetFPGA-public,” <https://github.com/NetFPGA/P4-NetFPGA-public/wiki>.